

Assignment 1: Classification

Key information

Deadlines

Submission: 11:59pm, 6 April, 2023 (Thursday week 7, Sydney time)

Late submissions policy

Late submissions are allowed for up to 3 days late. A penalty of 5% per day late will apply. Assignments more than 3 days late will not be accepted (i.e. will get 0 marks). The day cut-off time is 11:59pm.

Marking

This assignment is worth 15 marks = 15% of your final mark.

Your code will be marked for correctness. A few marks will be allocated for style – meaningful variable names and comments.

We will run your code. If it doesn't run, you will get 0 marks for the parts that don't run.

The assignment should be completed in pairs (groups of 2 students). No more than 2 students are allowed.

Submission

This assignment must be written in **Python** in the Jupyter Notebook environment. A Jupyter Notebook template is provided. Your implementation should use the same suite of libraries that we have used during the tutorials, such as **scikit-learn**, **numpy** and **pandas**.

The assignment will be submitted in Canvas.

Submission instructions:

- Before you submit, you need to create a group. In Canvas -> "People", select one of these two tabs: "A1part1" or "A1part2". Choose one of the empty groups listed and join it. Both you and your partner must join the same group. Groups have a maximum of 2 members.
- When you are ready to submit your assignment, you need to submit it on behalf of the group in the corresponding submission box. You need to submit **two versions of your code: ipynb and pdf**. Only one student from the group needs to submit, not both.
- In summary:
 - If you have registered your group under "**A1part1**", submit your ipynb code in "Submission: Assignment 1 ipynb for **A1part1** groups" and your pdf in "Submission: Assignment 1 pdf for **A1part1** groups"
 - If you have registered your group under "**A1part2**", submit your ipynb code in "Submission: Assignment 1 ipynb for **A1part2** groups" and your pdf in "Submission: Assignment 1 pdf for **A1part2** groups"

It is important to follow the submission instructions carefully as otherwise your mark may not be recorded correctly.

We had to create two options (A1part1 and A1part2) and two submission boxes because of the limitations of Canvas for the number of groups it allows.

File names and student names

- The submission files should be named like this: a1-SID1-SID2.ipynb (.pdf), where SID1 and SID2 are the SIDs of the two students
- In the Jupyter Notebook, include only your SIDs (as shown in the template) and not your name. The marking is anonymous.

Task

In this assignment you will investigate a real dataset by implementing multiple classification algorithms. You will first pre-process the dataset by replacing missing values and normalising the dataset with a min-max scaler. You will then evaluate the performance of multiple classification algorithms: K-Nearest Neighbour, Logistic Regression, Naïve Bayes, Decision Tree, Support Vector Machine, Bagging, AdaBoost, Gradient Boosting and Random Forest, using the stratified 10-fold cross-validation method. You will also apply a grid search to find the best parameters for some of these classifiers.

1. Data loading, pre-processing and printing

The dataset for this assignment is the Breast Cancer Wisconsin. It contains 699 examples described by 9 numeric attributes. There are two classes – **class1**, corresponding to benign breast cancer tumours, and **class2**, corresponding to malignant breast cancer tumours. The features are computed from a digitized image of a biopsy sample of breast tissue for a subject.

The dataset should be downloaded from Canvas: **breast-cancer-wisconsin.csv**. This file includes the attribute (feature) headings and each row corresponds to one individual. Missing attributes in the dataset are recorded with a ‘?’.

You will need to pre-process the dataset, before you can apply the classification algorithms. Three types of pre-processing are required: filling in the missing values, normalisation and changing the class values. After this is done, you need to print the first 10 rows of the pre-processed dataset.

1. Filling in the missing attribute values - The **missing attribute values** should be replaced with the mean value of the column using `sklearn.impute.SimpleImputer`.
2. Normalising the data - **Normalisation** of each attribute should be performed using a min-max scaler to normalise the values between [0,1] with `sklearn.preprocessing.MinMaxScaler`.
3. Changing the class values - The classes **class1** and **class2** should be changed to **0** and **1** respectively.
4. Print the first 10 rows of the pre-processed dataset. The feature values should be formatted to 4 decimal places using `.4f`, the class value is an integer. A function **print_data** has been provided in the template to help you achieve this.

For example, if your normalised data looks like this:

Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses	Class
0.1343	0.4333	0.5432	0.8589	0.3737	0.9485	0.4834	0.9456	0.4329	0
0.1345	0.4432	0.4567	0.4323	0.1111	0.3456	0.3213	0.8985	0.3456	1
0.4948	0.4798	0.2543	0.1876	0.9846	0.3345	0.4567	0.4983	0.2845	0

Then your program should print:

```
0.1343,0.4333,0.5432,0.8589,0.3737,0.9485,0.4834,0.9456,0.4329,0
0.1345,0.4432,0.4567,0.4323,0.1111,0.3456,0.3213,0.8985,0.3456,1
0.4948,0.4798,0.2543,0.1876,0.9846,0.3345,0.4567,0.4983,0.2845,0
```

(You need to print the first 10 rows not the first 3.)

Please note that we will test your code with another dataset, and your pre-processing should be written with this in mind. See the “Marking Criteria” section for more detail.

2. Defining functions for the classification algorithms

Part 1: Cross-validation without parameter tuning

You will now apply multiple classifiers to the pre-processed dataset, in particular: Nearest Neighbor, Logistic Regression, Naïve Bayes, Decision Tree, Bagging, Ada Boost and Gradient Boosting. All classifiers should use the **sklearn** modules from the tutorials. All random states in the classifiers should be set to **random_state=0**.

You need to evaluate the performance of these classifiers using 10-fold stratified cross-validation from **sklearn.model_selection.StratifiedKFold** with these options:

```
cvKFold=StratifiedKFold(n_splits=10, shuffle=True, random_state=0)
```

You will need to pass **cvKFold** (the stratified folds) as an argument when calculating the cross-validation accuracy, not **cv=10** as in the tutorials. This ensures that **random_state=0**.

For each classifier, write a function that accepts the required input and returns the average cross-validation score:

```
def exampleClassifier(X, y, [options]):
```

```
    ...
    return scores.mean()
```

where **X** contains the attribute values and **y** contains the class (as in the tutorial exercises).

More specifically, the headers of the functions for the classifiers are given below:

Logistic Regression

```
def logregClassifier(X, y)
```

```
    ...
    return scores.mean()
```

It should use **LogisticRegression** from **sklearn.linear_model**.

Naïve Bayes

```
def nbClassifier(X, y)
```

```
    ...
    return scores.mean()
```

It should use **GaussianNB** from **sklearn.naive_bayes**

Decision Tree

```
def dtClassifier(X, y)
```

```
    ...
    return scores.mean()
```

It should use DecisionTreeClassifier from sklearn.tree, with information gain (the entropy criterion)

Ensembles: Bagging, Ada Boost and Gradient Boosting

```
def bagDTClassifier(X, y, n_estimators, max_samples, max_depth)
```

```
    ...
    return scores.mean()
```

```
def adaDTClassifier(X, y, n_estimators, learning_rate, max_depth)
```

```
    ...
    return scores.mean()
```

```
def gbClassifier(X, y, n_estimators, learning_rate)
```

```
    ...
    return scores.mean()
```

These functions should implement Bagging, Ada Boost and Gradient Boosting using BaggingClassifier, AdaBoostClassifier and GradientBoostingClassifier from sklearn.ensemble. Bagging and Ada Boost should combine decision trees and use information gain.

Part 2: Cross-validation with parameter tuning

For two other classifiers, SVM and Random Forest, we would like to find the best parameters using grid search with 10-fold stratified cross-validation (GridSearchCV in sklearn).

The data should be split into training and test subsets using train_test_split from sklearn.model_selection with stratification and random_state=0 (as in the tutorials but with random_state=0).

You will need to pass cvKFold (the stratified folds) as an argument to GridSearchCV, not cv=10 as in the tutorials. This ensures that random_state=0.

Write the following functions:

K-Nearest Neighbour

```
def bestKNNClassifier(X, y)
```

```
    ...
    return (appropriate values so that the required printing can be done)
```

It should use the KNeighborsClassifier from sklearn.neighbors.

The grid search should consider the following values for the parameters $k(n_neighbors)$ and p :

$k = \{1, 3, 5, 7, 9\}$

$p = \{1, 2\}$

The function should return appropriate values, so that best parameters found, the best cross-validation accuracy and the test set accuracy can be printed when calling this function, see the next section.

SVM

def bestSVMClassifier(X,y)

...

return (appropriate values so that the required printing can be done)

It should use SVC from `sklearn.svm` with kernel set to 'rbf'.

The grid search should consider the following values for the parameters C and gamma:

$C = \{0.01, 0.1, 1, 5, 15\}$

$\text{gamma} = \{0.01, 0.1, 1, 10, 50\}$

The function should return appropriate values, so that best parameters found, the best cross-validation accuracy and the test set accuracy can be printed when calling this function, see the next section.

Random Forest

def bestRFClassifier(X,y)

It should use RandomForestClassifier from `sklearn.ensemble` with information gain and max_features set to 'sqrt'.

The grid search should consider the following values for the parameters n_estimators and max_leaf_nodes:

$n_estimators = \{10, 30, 60, 100, 150\}$

$max_leaf_nodes = \{6, 12, 18\}$

The function should return appropriate values, so that best parameters found, the best cross-validation accuracy, the test set accuracy, *the macro average F1 score, and the weighted average F1 score* can be printed when calling this function, see the next section.

3. Running the classifiers and printing the results

Run the classifiers from the previous section on the pre-processed dataset and print the results.

For Part 1 of this assignment, set the parameters as follows (this is already done for you in the template):

```
#Bagging
bag_n_estimators = 60
bag_max_samples = 100
bag_max_depth = 6

#AdaBoost
ada_n_estimators = 60
ada_learning_rate = 0.5
ada_bag_max_depth = 6
```

```
#GB
gb_n_estimators = 60
gb_learning_rate = 0.5
```

The printing should look like this but with the correct numbers (these are random numbers):

```
LR average cross-validation accuracy: 0.8123
NB average cross-validation accuracy: 0.7543
DT average cross-validation accuracy: 0.6345
Bagging average cross-validation accuracy: 0.8765
AdaBoost average cross-validation accuracy: 0.7165
GB average cross-validation accuracy: 0.9054

KNN best k: 7
KNN best p: 2
KNN cross-validation accuracy: 0.7853
KNN test set accuracy: 0.5991

SVM best C: 0.0100
SVM best gamma: 10.0000
SVM cross-validation accuracy: 0.8676
SVM test set accuracy: 0.8098

RF best n_estimators: 10
RF best max_leaf_nodes: 16
RF cross-validation accuracy: 0.8600
RF test set accuracy: 0.8321
RF test set macro average F1: 0.8123
RF test set weighted average F1: 0.8261
```

Format all numbers to 4 decimal places using .4f, except k, p, n_estimators and max_leaf_nodes which should be formatted as integers.

Academic honesty – very important

Please read the University policy on Academic Honesty very carefully:

<https://sydney.edu.au/students/academic-integrity.html>

Plagiarism (copying from another student, website or other sources), making your work available to another student to copy, engaging another person to complete the assignments instead of you (for payment or not) are all examples of **academic dishonesty**. Note that when there is copying between students, both students are penalised – the student who copies and the student who makes his/her work available for copying

The University penalties are severe and include: 1) a permanent record of academic dishonesty on your student file, 2) mark deduction, ranging from 0 for the assignment to Fail for the course and 3) expulsion from the University and cancelling of your student visa.

If there is a suspected case, the **investigation may take several months**. Your mark will not be finalised until the investigation is completed. This may create **problems enrolling in other courses next semester** (COMP5318 is a pre-requisite for many courses) or **delaying your graduation**. Going through the investigation is also **very stressful**.

In addition, the Australian Government passed a new legislation ([Prohibiting Academic Cheating Services Bill](#)) that makes it a **criminal offence** to provide or advertise academic cheating services - the provision or undertaking of work for students which forms a substantial part of a student's assessment task.

Do not confuse legitimate co-operation and cheating! You can discuss the assignment with other students but your group must write your own code.

We will use similarity detection software. If you cheat, the chances that you will be caught are very high.

Do not even think about engaging in plagiarism or academic dishonesty, it is not worth it. **Be smart and don't risk your future by engaging in plagiarism and academic dishonesty!**

Marking Criteria

The marking rubric is provided in Canvas.

Please note that we will test your program on another dataset. It will have the same format as the breast cancer dataset but a different number of features and examples, and different names of the features. You may assume the class value will be in the last column and there will be two classes as in the breast cancer dataset. The missing values may be everywhere, not only in a single column as the breast cancer dataset. Hence, do not hard-code the number of features and examples - do not set them to 699 and 9 as in the breast cancer dataset, and do not make assumptions that the missing values will be in a column with a specific name.

To test your code before submission, we have made available another dataset (**test-before.csv**), with the correct results.